中国人民大学
RENMIN UNIVERSITY OF CHINA
高瓴人工智能学院
Gaoling School of Artificial Intelligence

CoarsenConf: Equivariant Coarsening with Aggregated Attention for Molecular Conformer Generation

(ICLR 2024 Under Review)

Fanmeng Wang

2023-12-21

Outline

- Introduction
- Background
- Methods
- Experiment
- Conclusion

- Introduction
- Background
- Methods
- Experiment
- Conclusion

Introduction

- **Molecular conformer generation (MCG)** is a fundamental task in computational chemistry.
  - **Conformers** refer to the stable low-energy 3D molecular structures

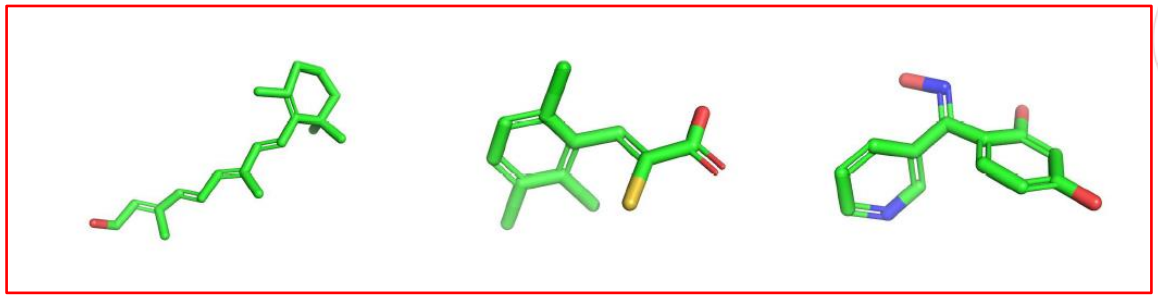

  - **Accurate molecular conformations** are important for various applications that depend on precise spatial and geometric qualities, including drug discovery and protein docking.

Introduction

- **Exciting Methods** for Molecular Conformation Generation
  - **Traditional physics-based methods**
    - a trade-off between speed and accuracy.
  - **Quantum mechanical methods**
    - more accurate but computationally slow.
  - **Stochastic cheminformatics-based methods** like RDKit ETKDG
    - more efficient but less accurate results.
  - **Existing generative MCG ML-based models**
    - fail to fully leverage the full geometric information inherent to the problem.

Introduction

- **CoarsenConf: an SE(3)-equivariant hierarchical VAE**
  - **Variable-length coarse-to-fine generation strategy**: coarse-grains molecules based on torsional angles, and create a flexible subgraph-level representation from corresponding fine-grained atom coordinates by an Aggregated Attention strategy.
  - CoarsenConf learns a **coarsegrained / subgraph-level latent distribution** for SE(3)-equivariant conformer generation.

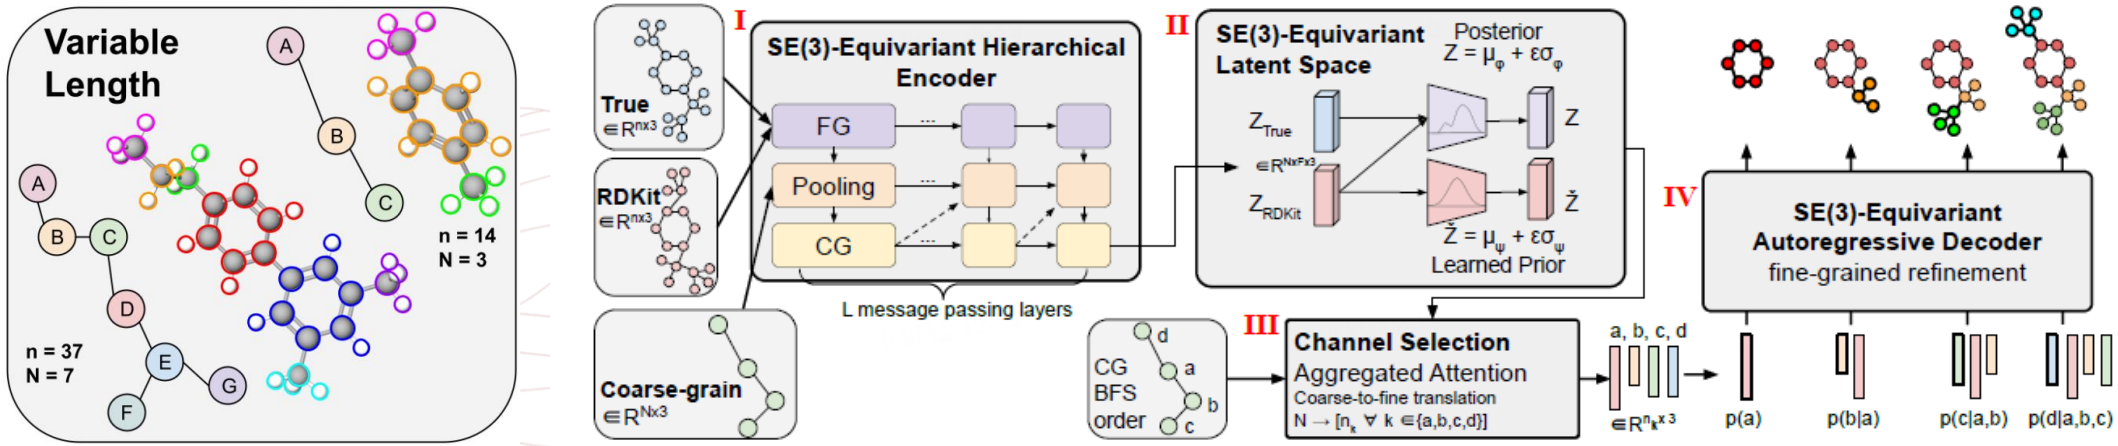

- Introduction
- Background
- Methods
- Experiment
- Conclusion

Background

■ Notations

- Each molecule is represented by a **graph** $G = (V, E)$ where V is the set of vertices representing atoms and E is the set of edges representing inter-atomic bonds.
- In addition, each molecular graph is expanded to incorporate **auxiliary edges** connecting all atoms within a 4Å radius to enhance long-range interactions in message passing.
- The **full molecular conformation** is represented by the coordinate matrix $\boldsymbol{X} \in \mathbb{R}^{|\mathcal{V}|\times 3}$

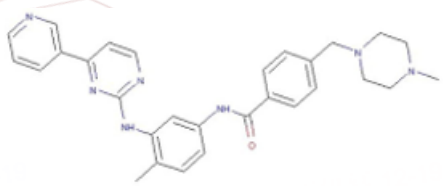

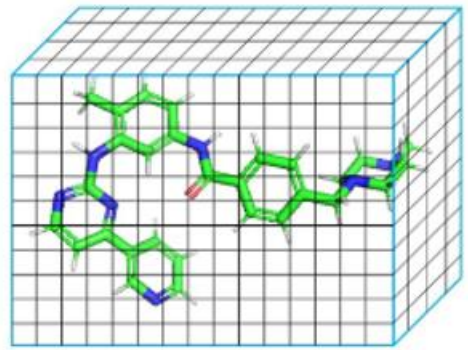

Molecular Graph Moelcular Conformation

Background

- **Problem Definition**
  - **Molecular conformation generation** is a conditional generative process that aims to model the conditional distribution of 3D molecular conformations $\boldsymbol{X}$, given the 2D molecule graph $\boldsymbol{G}$, i.e., $p(\boldsymbol{X}|G)$

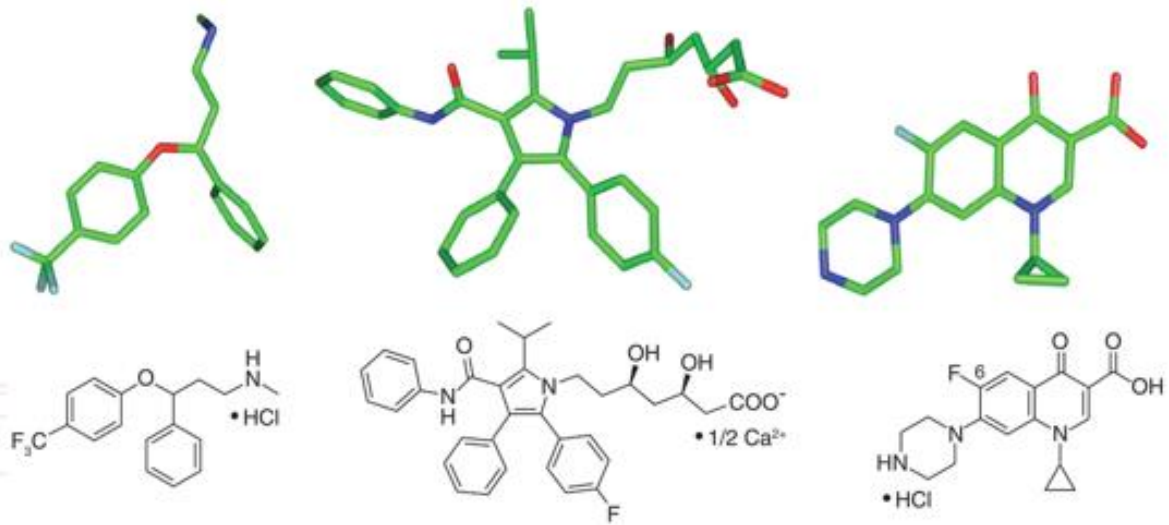

Background

■ Problem Definition

- Previous works have showed that RDKit is highly effective at generating conformations with **correct bond distances** and, as a result, can constrain the problem to a diffusion process over only torsion angles.
- In this paper, we thus formalize MCG as modeling the conditional distribution $p(\boldsymbol{X}|\mathcal{R})$ where $\mathcal{R}$ is the **RDKit generated atomic coordinates**, as we use RDKit as a building block to provide **an approximation starting** from only 2D information.

$$p(\boldsymbol{X}|G) \longrightarrow p(\boldsymbol{X}|\mathcal{R})$$

$$\mathcal{R} = \text{RDKit ETKDG}(G)$$

Background

■ Molecular Coarse-graining

- Molecular coarse-graining refers to the simplification of a molecule representation by grouping the **fine-grained (FG) atoms** in the original structure into **individual coarse-grained (CG) beads** with a rule-based mapping.

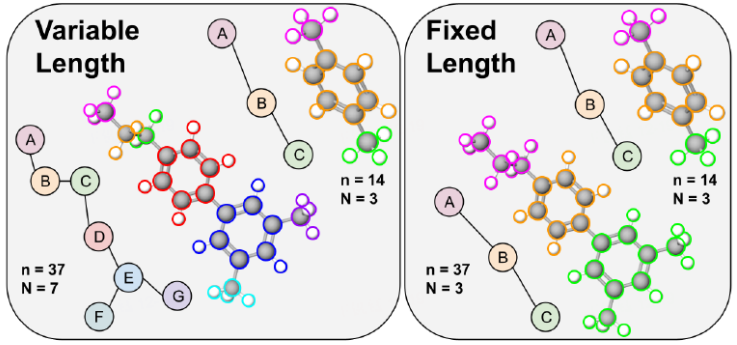

Figure 1: **Learning flexible coarse-grained representations.** CoarsenConf is the first model to employ variable-length (on left) coarse-graining. Each input molecule (n fine-grained (FG) atoms) can be represented by a different number of coarse-grained (CG) nodes N, thus accommodating diverse molecular sizes. In contrast, prior approaches rely on fixed-length (on right) coarse-graining, thereby forcing all molecules to possess the same number of CG nodes. Variable-length coarse-graining enhances the model's ability to create better learned representations across molecules of different sizes and geometries. The molecules on the left are coarsened along torsional angles.

- In this work, we coarse-grains the given molecule based on torsional angles, thus creating a **flexible variable-length CG representation**.

- Introduction
- Background
- Methods
- Experiment
- Conclusion

Method

- CoarsenConf is a **conditional generative model** that learns $p(X|R)$ where X is the **low-energy 3D conformation**, and R is the **RDKit approximate conformation**.

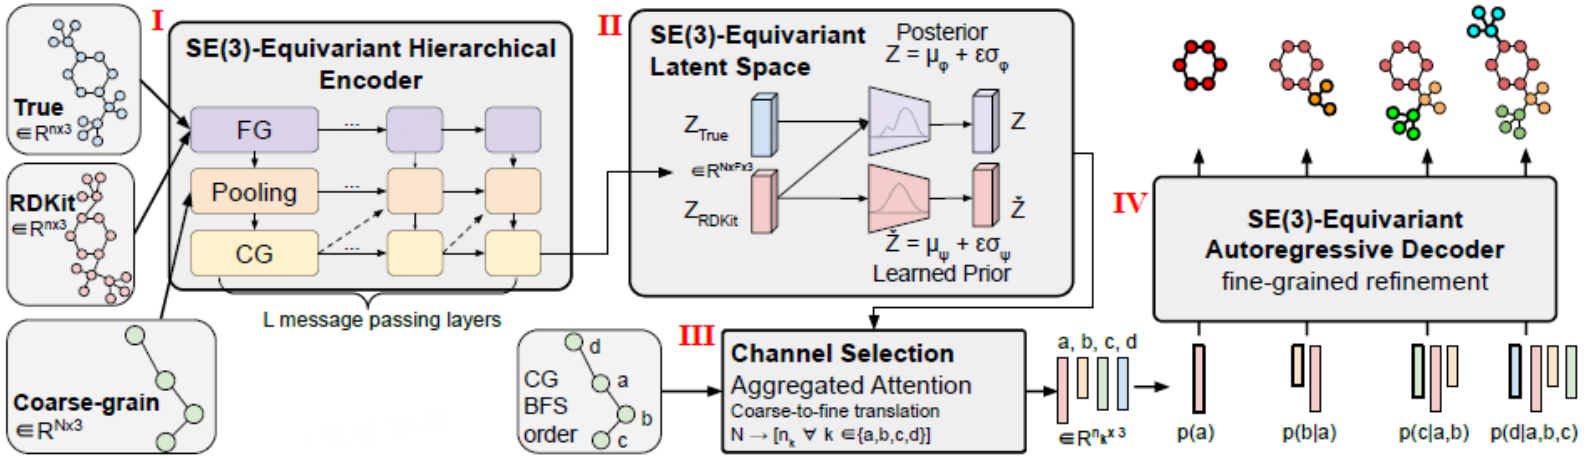

An SE(3)-equivariant hierarchical VAE architecture

Method: 3D Infomax

■ Encoder Architecture

➢ The encoder module takes in **SE(3)-equivariant atomistic coordinates** $x \in R^{n\times 3}$ and **SE(3)-invariant atom features** $h \in R^{n\times D}$ from the ground truth and RDKit conformer and creates coarse-grained latent representations for each Z and $\tilde{Z} \in R^{N\times F\times 3}$ where N is the number of CG beads, and F is the latent dimensions.

➢ A single encoder layer is composed of three modules: fine-grained, pooling, and coarse-grained.

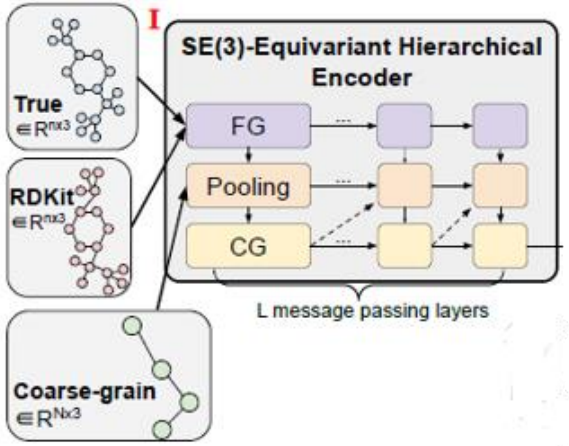

Method: 3D Infomax

■ Encoder Architecture

Algorithm 1 Encoder Forward Pass: Hierarchical Message Passing Inputs and Outputs

true coord X, RDKit coord $\hat{X}$
pooing coord X_p, pooling RDKit coord $\hat{X}_p$
coarse coord X_c, coarse RDKit coord $\hat{X}_c$
true features h, RDKit features $\hat{h}$
pooing features h_p, pooling RDKit features $\hat{h}_p$
coarse features h_c, coarse RDKit features $\hat{h}_c$
true latent CG representation Z, RDKit latent CG representation $\hat{Z}$
$(FG), (PL), (CG) \leftarrow$ dataloader$[i]$ *// Fine-grain, Pooling, and Coarse-grain graphs*
$(X, h), (\hat{X}, \hat{h}) \leftarrow FG$ *//* $X \in \mathbb{R}^{n\times 3}$ *and* $h \in \mathbb{R}^{n\times D}$
$(X_p, h_p), (\hat{X}_p, \hat{h}_p) \leftarrow PL$ *//* $X_p \in \mathbb{R}^{n+N\times 3}$ *and* $h_p \in \mathbb{R}^{n+N\times D}$
$(X_c, h_c), (\hat{X}_c, \hat{h}_c) \leftarrow CG$ *//* $X_c \in \mathbb{R}^{N\times 3}$ *and* $h_c \in \mathbb{R}^{N\times D}$
$Z, \hat{Z} \leftarrow 0, 0$ *//* $Z = [v_I \; \forall I \in X_c]$*, init as zeros* $\in \mathbb{R}^{N\times F\times 3}$*, see Eq. 8*
for t in num_layers **do**
  $(X, h), (\hat{X}, \hat{h}) \leftarrow$ FG_Module$((X, h), (\hat{X}, \hat{h}))$ *// see Eq. 6*
  $X_p[0:n] \leftarrow X$
  $\hat{X}_p[0:n] \leftarrow \hat{X}$
  $h_p[0:n] \leftarrow h$
  $\hat{h}_p[0:n] \leftarrow \hat{h}$ *// Set pooling graphs features with output of FG Module*
  $X_p[n:n+N] \leftarrow X_c$
  $\hat{X}_p[n:n+N] \leftarrow \hat{X}_c$
  $h_p[n:n+N] \leftarrow h_c$
  $\hat{h}_p[n:n+N] \leftarrow \hat{h}_c$ *// Set pooling graphs features with output of CG Module*
  $(X_p, h_p), (\hat{X}_p, \hat{h}_p) \leftarrow$ Pooling_Module$((X_p, h_p), (\hat{X}_p, \hat{h}_p))$ *// see Eq. 7*
  $X_c \leftarrow X_p[n:n+N]$
  $\hat{X}_c \leftarrow \hat{X}_p[n:n+N]$
  $h_c \leftarrow h_p[n:n+N]$
  $\hat{h}_c \leftarrow \hat{h}_p[n:n+N]$ *// Set CG graphs features with output of Pooling Module*
  $Z, \hat{Z} \leftarrow$ CG_Module$((X_c, h_c), (\hat{X}_c, \hat{h}_c), Z, \hat{Z})$ *// [see Eq. 8 & Eq. 9]*
end for
Return $Z, \hat{Z} \in \mathbb{R}^{N\times F\times 3}$

FG Module

$$
\begin{aligned}
\boldsymbol{m}_{j\to i} &= \phi^e(\boldsymbol{h}_i^{(t)}, \boldsymbol{h}_j^{(t)}, \|\boldsymbol{x}_i^{(t)} - \boldsymbol{x}_j^{(t)}\|^2, \boldsymbol{f}_{j\to i}), \forall (I, J) \in \mathcal{E} \cup \mathcal{E}', \\
\boldsymbol{u}_{j'\to i} &= a_{j'\to i} \boldsymbol{W} \boldsymbol{h}_{j'}^{(t)}, \forall i \in \mathcal{V}, j' \in \mathcal{V}', \\
\boldsymbol{m}_i &= \frac{1}{|\mathcal{N}(i)|} \sum_{j\in\mathcal{N}(i)} \boldsymbol{m}_{j\to i}, \forall i \in \mathcal{V} \cup \mathcal{V}', \\
\boldsymbol{u}_i &= \sum_{j'\in\mathcal{V}'} \boldsymbol{u}_{j'\to i}, \forall i \in \mathcal{V}, \quad \text{and} \quad \boldsymbol{u}_i' = 0, \\
\boldsymbol{x}_i^{(t+1)} &= \eta_x \cdot \boldsymbol{x}_i^{(0)} + (1 - \eta_x) \cdot \boldsymbol{x}_i^{(t)} + \sum_{j\in\mathcal{N}(i)} (\boldsymbol{x}_i^{(t)} - \boldsymbol{x}_j^{(t)}) \phi^x(\boldsymbol{m}_{j\to i}), \\
\boldsymbol{h}_i^{(t+1)} &= (1 - \eta_h) \cdot \boldsymbol{h}_i^{(t)} + \eta_h \cdot \phi^h(\boldsymbol{h}_i^{(t)}, \boldsymbol{m}_i, \boldsymbol{u}_i, \boldsymbol{f}_i), \forall i \in \mathcal{V} \cup \mathcal{V}',
\end{aligned}
\tag{6}
$$

Pooling Module

$$
\begin{aligned}
\boldsymbol{m}_{j\to I} &= \phi^e(\boldsymbol{H}_I^{(t)}, \boldsymbol{h}_j^{(t)}, \|\boldsymbol{X}_I^{(t)} - \boldsymbol{x}_j^{(t)}\|^2, \boldsymbol{f}_{j\to I}), \forall (I, J) \in \mathcal{E} \cup \mathcal{E}', \\
\boldsymbol{m}_I &= \frac{1}{|\mathcal{N}(I)|} \sum_{j\in\mathcal{N}(I)} \boldsymbol{m}_{j\to I}, \forall I \in \mathcal{V} \cup \mathcal{V}', \\
\boldsymbol{X}_I^{(t+1)} &= \eta_X \cdot \boldsymbol{X}_I^{(0)} + (1 - \eta_X) \cdot \boldsymbol{X}_I^{(t)} + \sum_{j\in\mathcal{N}(I)} (\boldsymbol{X}_I^{(t)} - \boldsymbol{x}_j^{(t)}) \phi^x(\boldsymbol{m}_{j\to I}), \\
\boldsymbol{H}_I^{(t+1)} &= (1 - \eta_H) \cdot \boldsymbol{H}_I^{(t)} + \eta_H \cdot \phi^h(\boldsymbol{H}_I^{(t)}, \boldsymbol{m}_I, \boldsymbol{f}_I), \forall I \in \mathcal{V} \cup \mathcal{V}',
\end{aligned}
\tag{7}
$$

CG Module

$$H_I' = \phi_1(h_I^{(t)}, \|\text{VN-MLP}_1(\boldsymbol{v}_I^{(t)})\|) \in \mathbb{R}^D, \tag{8a}$$

$$H_I'' = \phi_2(h_I^{(t)}, \|\text{VN-MLP}_2(\boldsymbol{v}_I^{(t)})\|) \in \mathbb{R}^F, \tag{8b}$$

$$\boldsymbol{v}_I' = \text{diag}\{\phi_3(H_I^{(t)})\} \cdot \text{VN-MLP}_3(\boldsymbol{v}_I^{(t)}) \in \mathbb{R}^{F\times 3}. \tag{8c}$$

$$\boldsymbol{m}_{I\leftarrow J}^{\boldsymbol{H}} = \text{Ker}_1(\|\boldsymbol{r}_{I,J}\|) \odot H_J', \tag{9a}$$

$$\boldsymbol{m}_{I\leftarrow J}^{\boldsymbol{v}} = \text{diag}\{\text{Ker}_2(\|\boldsymbol{r}_{I,J}\|)\} \cdot \boldsymbol{v}_J' + (\text{Ker}_3(\|\boldsymbol{r}_{I,J}\|) \odot H_J'') \cdot \boldsymbol{r}_{I,J}^\top, \tag{9b}$$

$$\boldsymbol{u}_{J'\to I} = a_{J'\to I} \boldsymbol{W} \boldsymbol{H}_{J'}^{(t)}, \forall I \in \mathcal{V}, J' \in \mathcal{V}', \tag{9c}$$

$$\boldsymbol{u}_I = \sum_{J'\in\mathcal{V}'} \boldsymbol{u}_{J'\to I}, \forall I \in \mathcal{V}, \quad \text{and} \quad \boldsymbol{u}_I' = 0, \tag{9d}$$

$$\boldsymbol{H}_I^{t+1} = (1 - \eta_H) \cdot H_I^\ell + \eta_H \cdot \text{MLP}(H_I^\ell, \sum_{J\in N(I)} m_{I\leftarrow J}^H, u_I), \forall I \in \mathcal{V} \cup \mathcal{V}', \tag{9e}$$

$$\boldsymbol{v}_I^{t+1} = (1 - \eta_v) \cdot \boldsymbol{v}_I^\ell + \eta_v \cdot \text{VN-MLP}_4(\boldsymbol{v}_I^\ell, \sum_{J\in N(I)} \boldsymbol{m}_{I\leftarrow J}^{\boldsymbol{v}}), \forall I \in \mathcal{V} \cup \mathcal{V}', \tag{9f}$$

Method: 3D Infomax

■ Equivariant Latent Space

➢ The **conditional posterior** is parameterized with both the ground truth and RDKit approximation, whereas the **learned conditional** prior only uses the RDKit.

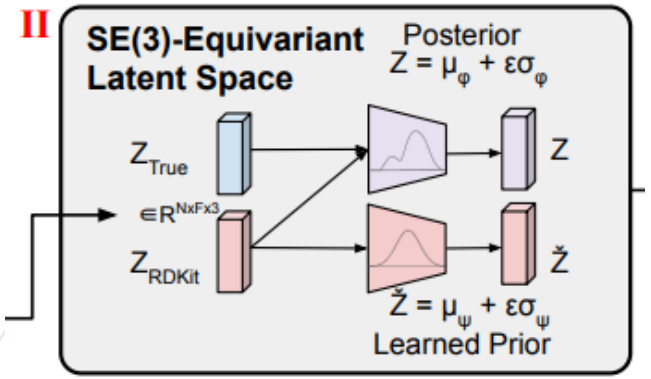

$$\text{Posterior}: \boldsymbol{\mu}_\phi = \text{VN-MLP}(Z, \tilde{Z}), \quad \log(\boldsymbol{\sigma}_\phi^2) = \text{MLP}(Z, \tilde{Z}),$$
$$\text{Prior}: \boldsymbol{\mu}_\psi = \text{VN-MLP}(\tilde{Z}), \quad \log(\boldsymbol{\sigma}_\psi^2) = \text{MLP}(\tilde{Z}). \tag{2}$$

Method: 3D Infomax

■ Decoder Architecture: Channel Selection

➢ Given latent representation Z is still in CG space, we need to perform variable-length backmapping to convert back to FG space so that we can further refine the atom coordinates to generate the low energy conformer

$$X \leftarrow \text{Aggregated\_Attention}(\text{KV} = Z, \text{Q} = \hat{X})$$

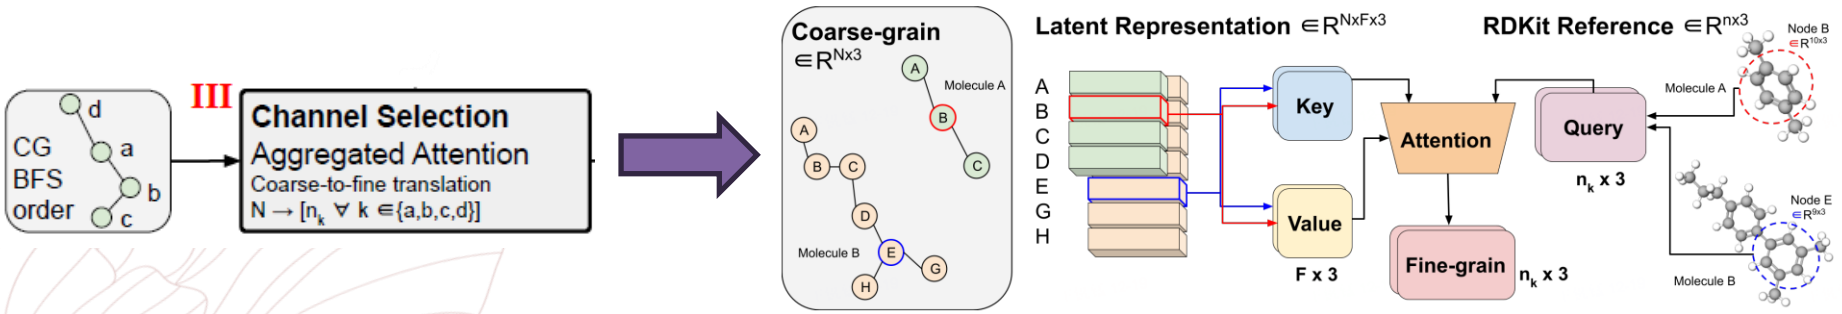

Figure 3: **Variable-length coarse-to-fine backmapping via Aggregated Attention.** The highlighted latent beads of two independent molecules are attended to by the respective fine-grained queries in a batched manner (see red and blue), to generate FG coordinates in the desired shape (matching input queries on right). The single-head attention operation uses the latent vectors of each CG bead $Z_k \in R^{F\times 3}$ for each molecule as the keys and values, with an embedding dimension of 3 to match the x, y, z coordinates. The query vectors are the FG subset of the respective RDKit conformers, corresponding to each CG bead $\in R^{n_k\times 3}$. We know a priori how many FG atoms correspond to a certain CG bead (n_k). Aggregated Attention learns the optimal blending of CG features for FG reconstruction by aggregating 3D segments of FG information to form our latent query.

Method: 3D Infomax

■ Decoder Architecture: Channel Selection

- The decoder architecture is similar to the EGNN-based FG module in the encoder, but we just learn to predict the **difference** between the RDKit conformations and ground truth conformations.
- We simplify the learning objective by setting X = R + ΔX, and learn the optimal distortion ΔX from the RDKit approximation.

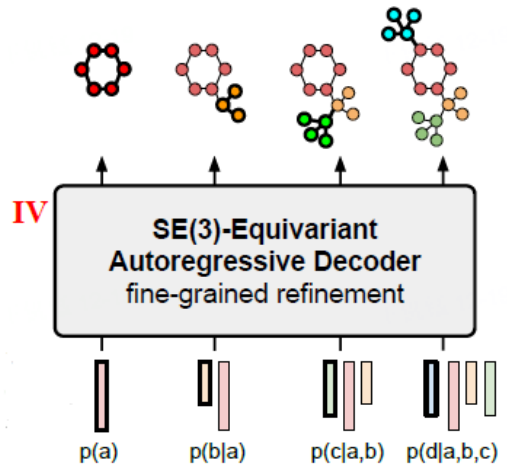

More formally, a single decoder layer is defined as follows:

$$\boldsymbol{\mu}^{(t)} = \frac{1}{|\mathcal{V}_{prev}|} \sum_{k \in \mathcal{V}_{prev}} x_k, \tag{10a}$$

$$\tilde{\boldsymbol{h}}_i = \phi^m(\boldsymbol{h}_i^{(t)}, \boldsymbol{x}_i^{(t)}, \boldsymbol{\mu}^{(t)}, \|\boldsymbol{x}_i^{(t)} - \boldsymbol{\mu}^{(t)}\|^2), \forall i \in \mathcal{V}_{cur}, \tag{10b}$$

$$\boldsymbol{m}_{j \to i} = \phi^e(\tilde{\boldsymbol{h}}_i^{(t)}, \tilde{\boldsymbol{h}}_j^{(t)}, \|\boldsymbol{x}_i^{(t)} - \boldsymbol{x}_j^{(t)}\|^2, \|\boldsymbol{x}_i^{(t)} - \boldsymbol{x}_{ref,j}^{(t)}\|^2, \|\boldsymbol{x}_i^{(t)} - \boldsymbol{x}_{ref,i}^{(t)}\|^2), \forall (i,j) \in \mathcal{E}_{cur}, \tag{10c}$$

$$\boldsymbol{m}_i = \frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} \boldsymbol{m}_{j \to i}, \forall i \in \mathcal{V}_{cur}, \tag{10d}$$

$$\boldsymbol{u}_{j' \to i} = a_{j' \to i} \boldsymbol{W} \boldsymbol{h}_{j'}^{(t)}, \forall i \in \mathcal{V}_{cur}, j' \in \mathcal{V}_{prev}, \tag{10e}$$

$$\boldsymbol{u}_i = \sum_{j' \in \mathcal{V}_{prev}} \boldsymbol{u}_{j' \to i}, \forall i \in \mathcal{V}_{cur}, \tag{10f}$$

$$\boldsymbol{x}_i^{(t+1)} = \boldsymbol{x}_{ref,i}^{(t)} + \sum_{j \in \mathcal{N}(i)} (\boldsymbol{x}_i^{(t)} - \boldsymbol{x}_j^{(t)}) \phi^x(\boldsymbol{m}_{j \to i}), \forall i \in \mathcal{V}_{cur}, \tag{10g}$$

$$\boldsymbol{h}_i^{(t+1)} = (1 - \beta) \cdot \boldsymbol{h}_i^{(t)} + \beta \cdot \phi^h(\tilde{\boldsymbol{h}}_i^{(t)}, \boldsymbol{m}_i, \boldsymbol{u}_i, \boldsymbol{f}_i), \forall i \in \mathcal{V}_{cur}, \tag{10h}$$

Method: 3D Infomax

General algorithms

Algorithm 2 Training One Epoch

for data in dataloader **do**
  $(FG), (PL), (CG) \leftarrow$ data
  $(X, h), (\hat{X}, \hat{h}) \leftarrow FG$
  $Z, \hat{Z} \leftarrow$ Encoder$(data)$ *// Refer to Algorithm 1*
  $Z \leftarrow$ Posterior$(Z, \hat{Z})$ *// see Eq. 2*
  $\hat{Z} \leftarrow$ Prior$(\hat{Z})$ *// see Eq. 2*
  $X \leftarrow$ Aggregated_Attention(KV $= Z$, Q $= \hat{X}$)
  $X_{Gen} \leftarrow$ Decoder(X)
  $loss \leftarrow$ Loss$(X, X_{\text{Gen}}) + \text{KL}(Z, \hat{Z})$ *// see Eq. 5*
  $loss$.backward()
end for

Algorithm 3 Inference

for data in dataloader **do**
  $(FG), (PL), (CG) \leftarrow$ data
  $(X, h), (\hat{X}, \hat{h}) \leftarrow FG$
  $FG \leftarrow ((\hat{X}, \hat{h}), (\hat{X}, \hat{h}))$ *// use RDKit for both, discard true structure in inference*
  $PL \leftarrow (PL[1], PL[1])$ *// discard true structure in inference*
  $CG \leftarrow (CG[1], CG[1])$ *// discard true structure in inference*
  None, $\hat{Z} \leftarrow$ Encoder((FG, PL, CG)) *// Refer to Algorithm 1*
  $Z \leftarrow$ Prior$(\hat{Z})$ *// see Eq. 2*
  $X \leftarrow$ Aggregated_Attention(KV $= Z$, Q $= \hat{X}$)
  $X_{Gen} \leftarrow$ Decoder(X)
  results $\leftarrow X_{Gen}$
end for
return results

Loss function (one-to-one loss)

$$\text{MSE}(\mathcal{A}(X, X_{true})) + \beta_1 D_{KL}(q_\phi(z|X, \mathcal{R}) \parallel p_\psi(z|\mathcal{R})) + \beta_2 \frac{1}{|\mathcal{E}^*|} \sum_{(i,j)\in\mathcal{E}^*} ||r_{ij} - r_{ij}^{true}||^2, \quad (5)$$

Optimal Transport reduces compute requirements

$$\mathcal{L}_{OT} = \min_{\mathbf{T}\in\mathcal{Q}_{K,L}} \sum_{k,l} T_{kl} \mathcal{L}(\mathcal{C}_k, \mathcal{C}_l^*),$$

$$\mathcal{L}(\mathcal{C}_k, \mathcal{C}_l^*) = \text{MSE}(\mathcal{C}_k, \mathcal{C}_l^*) + \text{distance error}(\mathcal{C}_k, \mathcal{C}_l^*),$$

- Introduction
- Background
- Methods
- **Experiment**
- Conclusion

Experiment

■ GEOM BENCHMARKS: 3D COORDINATE RMSD

➢ **GEOM dataset:** consisting of QM9 (average 11 atoms, 15 conformers per molecule) and DRUGS (average 44 atoms, 104 per molecule)

Table 1: Quality of ML generated conformer ensembles for the GEOM-QM9 ($\delta = 0.5$Å) and GEOM-DRUGS ($\delta = 0.75$Å) test set in terms of Coverage (%) and Average RMSD (Å) Precision. Bolded results are the best, and the underlined results are second best. See Appendix §I- §J for more details.

	QM9-Precision				DRUGS-Precision			
	Coverage ↑		AMR ↓		Coverage ↑		AMR ↓	
Method	Mean	Med	Mean	Med	Mean	Med	Mean	Med
GeoDiff	50.0	33.5	0.524	0.510	23.7	13.0	1.131	1.083
GeoMol	75.9	**100.0**	0.262	0.233	40.5	33.5	0.919	0.842
Torsional Diffusion ($\ell = 2$)	<u>78.4</u>	**100.0**	<u>0.222</u>	<u>0.197</u>	**52.1**	**53.7**	**0.770**	0.720
CoarsenConf-OT	**80.2**	**100.0**	**0.149**	**0.107**	<u>52.0</u>	<u>52.1</u>	<u>0.836</u>	**0.694**

➢ Metrics:

$$\text{COV-Precision} := \frac{1}{K}\left|\{k \in [1..K] : \min_{l\in[1..L]} \text{RMSD}(C_k, C_l^*) < \delta\}\right|,$$

$$\text{AMR-Precision} := \frac{1}{K}\sum_{k\in[1..K]} \min_{l\in[1..L]} \text{RMSD}(C_k, C_l^*),$$

➢ We outperform all models on QM9, and yield competitive results on DRUGS when using an optimal transport (OT) loss.

Experiment

■ GEOM BENCHMARKS: PROPERTY PREDICTION

Table 2: Property prediction: Mean absolute error of generated vs. ground truth ensemble properties for E, HOMO-LUMO gap $\Delta\epsilon$, E_{min} (kcal/mol), and dipole moment μ (debye) calculated with xTB.

	E	μ	$\Delta\epsilon$	E_{min}
DMCG	-	-	-	0.136
GeoDiff	-	-	-	0.155
GeoMol	28.80	1.475	4.186	0.267
Torsional Diffusion	16.75	1.333	2.908	0.096
CoarsenConf	**12.41**	**1.250**	**2.522**	**0.049**

- CoarsenConf is able to generate **the lowest energy structures** with **the most accurate chemical properties**.

Experiment

Via Autodock Vina's flexible docking simulation

■ FLEXIBLE ORACLE-BASED PROTEIN DOCKING

- We evaluate MCG models, pretrained on GEOM-DRUGS, using nine protein docking oracle functions provided by the Therapeutics Data Commons (TDC)
- For each evaluated MCG method, we generate 50 conformers for each of the nine ligands and report the best (lowest) binding affinity.

Table 3: Quality of best generated conformer for known protein ligands for all 9 proteins from the TDC library. Quality is measured by free energy change (kcal/mol) of the binding process with AutoDock Vina's flexible docking simulation (↓ is better).

	Best Protein-Conformer Binding Affinity (↓ is better)								
Method	3PBL	2RGP	1IEP	3EML	3NY8	4RLU	4UNN	5M04	7L11
RDKit + MMFF	-8.26	-11.42	-10.75	-9.26	-9.69	-8.72	-9.73	-9.53	-9.19
GeoMol	-8.23	-11.49	-11.16	-9.39	**-11.66**	-8.85	-10.28	-9.31	-9.29
Torsional Diffusion	-8.53	-11.34	-10.76	-9.25	-10.32	-8.96	-10.65	-9.61	-9.10
CoarsenConf	**-8.81**	**-12.93**	**-16.43**	**-9.82**	-11.26	**-9.54**	**-11.62**	**-14.00**	**-9.43**

- CoarsenConf significantly outperforms prior MCG methods on the TDC oracle-based affinity prediction task.

- Introduction
- Background
- Methods
- **Experiment**
- Conclusion

Conclusion

- We present CoarsenConf, a novel approach for robust molecular conformer generation that combines an SE(3)-equivariant hierarchical VAE with geometric coarse-graining strategy for accurate conformer generation.
- By utilizing easy-to-obtain approximate conformations generated by RDKit, our model effectively learns to generate low-energy conformers.
- Our experiments demonstrate the effectiveness of CoarsenConf compared to existing methods.

Thank You for listening!

Fanmeng Wang

2023-12-21